

Persistence

Programming 4

Agenda




1. Persistence stores
2. Repositories
3. Datamodels
4. Transactions
5. Object-Relational Mapping

Persistence stores

Definition

Persistence:
state that outlives
the process that created it.

lives longer than your process (in memory)



Persistence stores

data access obj (repository) change
in between strategies to save time


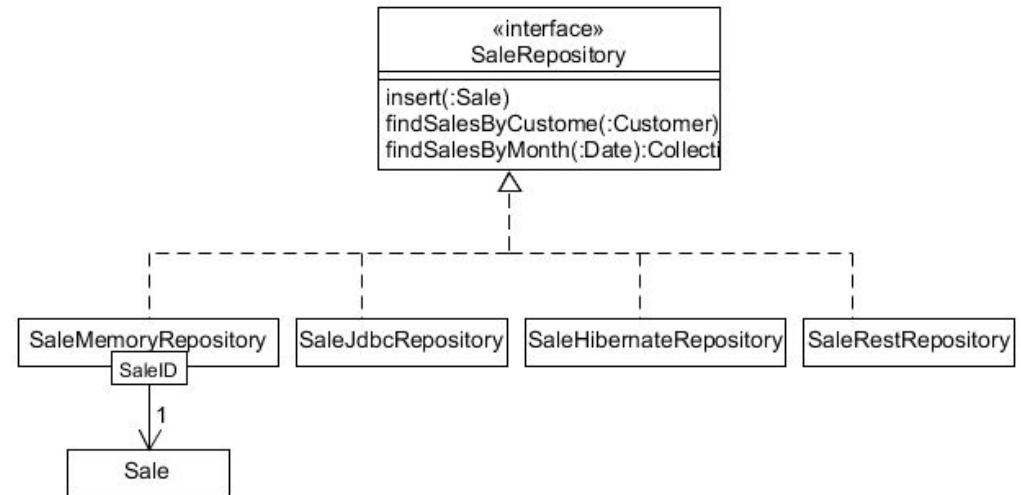
❑ Structured data



❑ Databases



specialized PS



Structured data: Marshalling

serialization

- In memory data have any graph structure
 - Pointers in all directions
- Persistence requires conversion to a byte stream (marshalling) and back (unmarshalling)
 - Send to file (persistence)
 - Only one process can access the file at a time
 - Send to network (Remote communication)
- Storing data that can be read by another program
 - Not transactional

only one project can access a file at a time

Java serialization

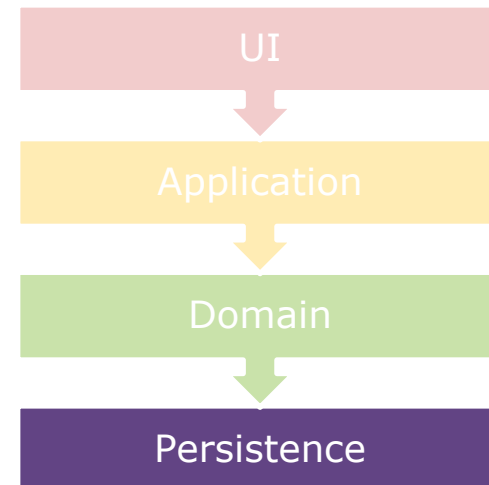

- ❑ Builtin API that automatically converts objects to bytestream
- ❑ Binary
- ❑ Supports graphs
- ❑ Java specific, OS/hardware independent
 - ❑ For deserialising you need to have all class code available in receiving application
- ❑ Usage
 - ❑ Save small amounts of data
 - ❑ Remote Method Invocation

java specific format

save class info but not the class itself

the other program needs to know the classes so that it can deserialize it

serialization is recursive: saves obj and all of their attributes



Java serialization

Serializable interface required
Marker (empty) interface

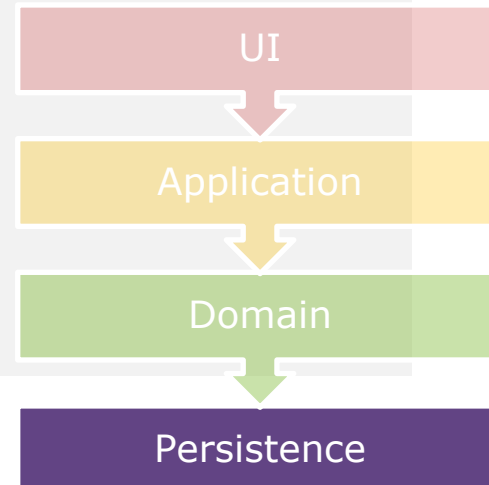
SerializeDemo.java

```
public class Student implements Serializable{
    private final int studentIdentity;
    private final String name;
    private final LocalDate birthDate;
    private transient String residence;
    //...
}

public class SerializeDemo {
    public static void main(String[] args) {
        List<Student> studentList = List.of(
            new Student(9999, "Charlotte Vermeulen", LocalDate.of(2000, 1, 24), "Antwerp"),
            new Student(666, "Donald Muylle", LocalDate.of(1952, 6, 14), "Roeselare"));
        try (ObjectOutputStream oos = new ObjectOutputStream(
            new FileOutputStream("data/students.ser"))) {
            oos.writeObject(studentList);
        } catch (IOException e) { e.printStackTrace();}
        // ...
    } // end main
} // end SerializeDemo
```

if you don't implement the Serializable interface
you'll not be able to serialize anything that wasn't
specifically marked as serializable


Transient fields are not serialized



Marshalling to structured text

- ❑ **Saving limited amounts of language independent data**

```
<?xml version="1.0"?>
<quiz>
  <qanda seq="1">
    <question>
      Who was the forty-second
      president of the U.S.A.?
    </question>
    <answer>
      William Jefferson Clinton
    </answer>
  </qanda>
  <!-- Note: We need to add
  more questions later.-->
</quiz>
```



- ❑ **CSV**
 - ❑ Not a strict format, requires additional conventions

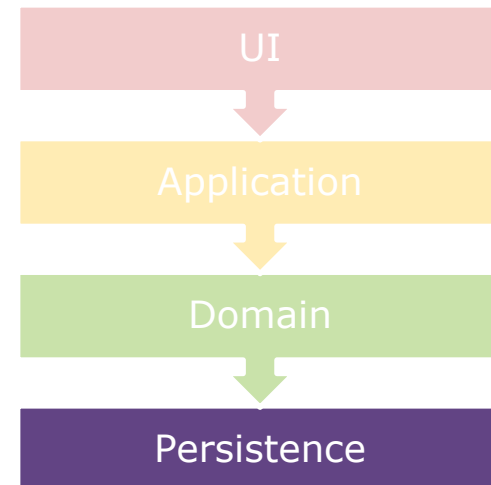
{ JSON }
JavaScript Object Notation

language and platform independent

- ❑ **json, xml, yaml...** ASCII formats
 - ❑ Tree structures: need a strategy to handle graphs
 - ❑ Used in webservices (REST)
 - ❑ Libraries: JAXB, Jackson, Moshi

they work in between diff languages

jason b: binding java objects to json



Databases

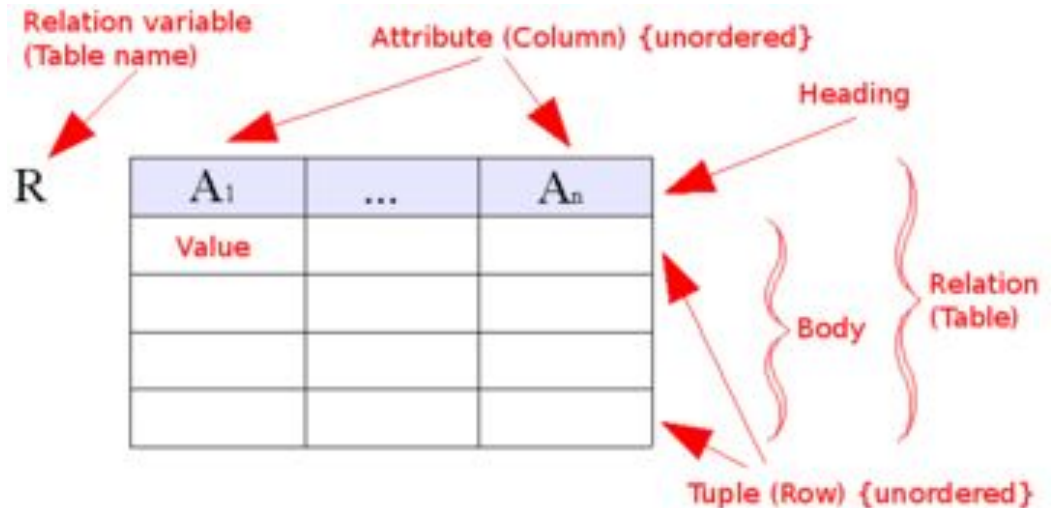
- ❑ Parallel access simultaneous access by multiple systems
- ❑ Client-server
- ❑ Transactional can't save and retreat data in step
- ❑ large amounts of data more persistent technologies



Relational databases

- ❑ Most common type of database
other types gain importance(noSQL)
- ❑ Predefined datatypes and structure
database schema with multiple tables
- ❑ Tables contain unique records (= tuples, rows)
primary key
- ❑ Relations using keys
foreign keys

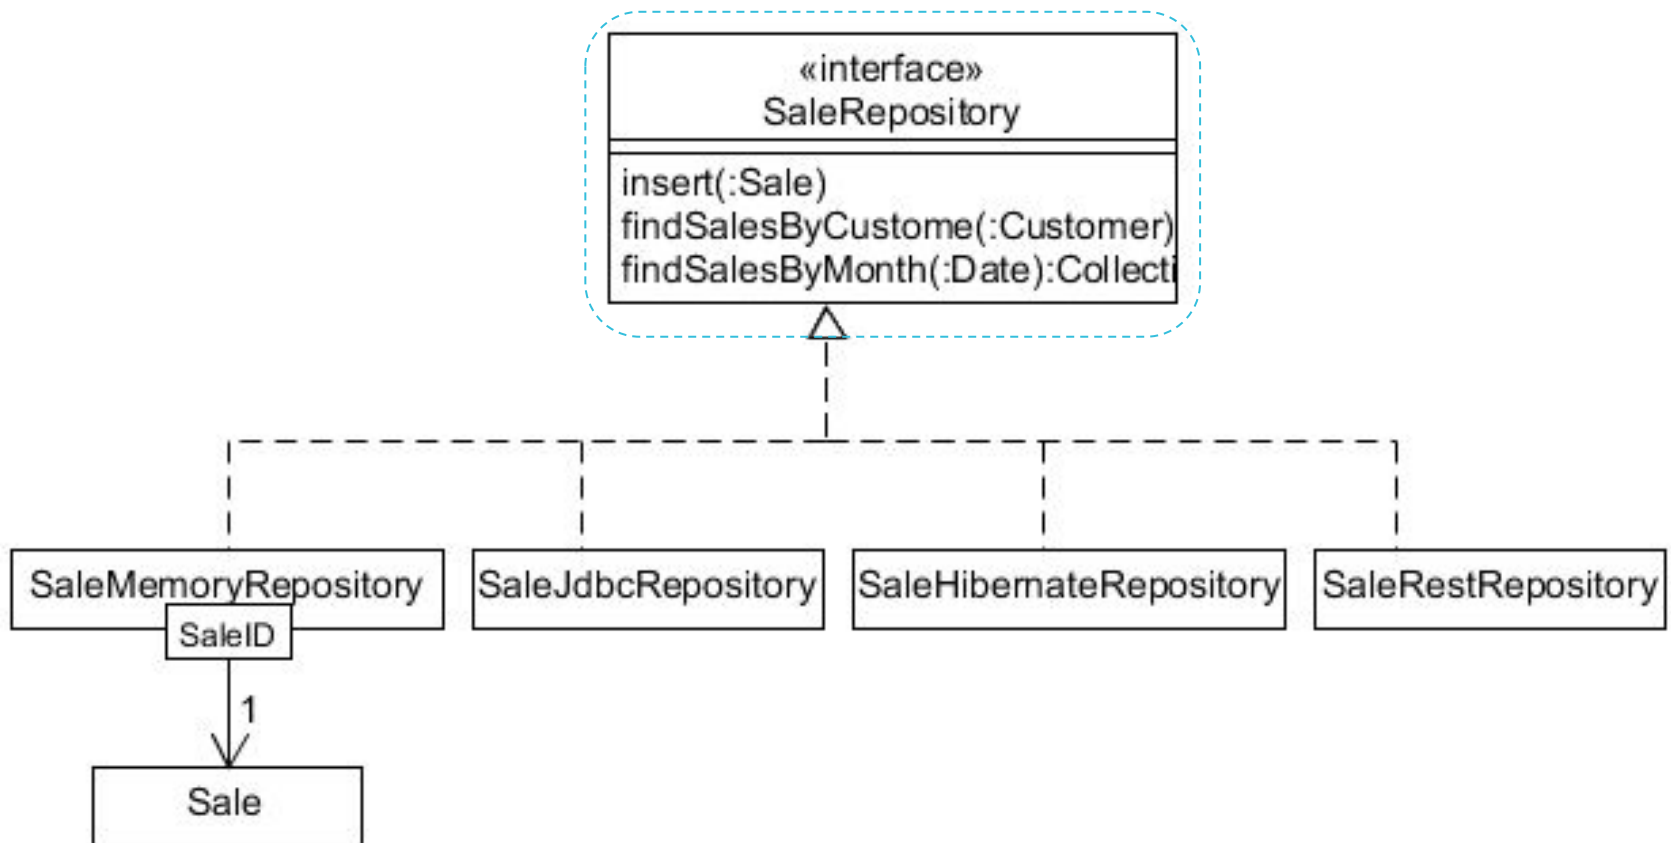
json = open api(documenting what the json structure is for and if they're compliant for the schema)



Repositories

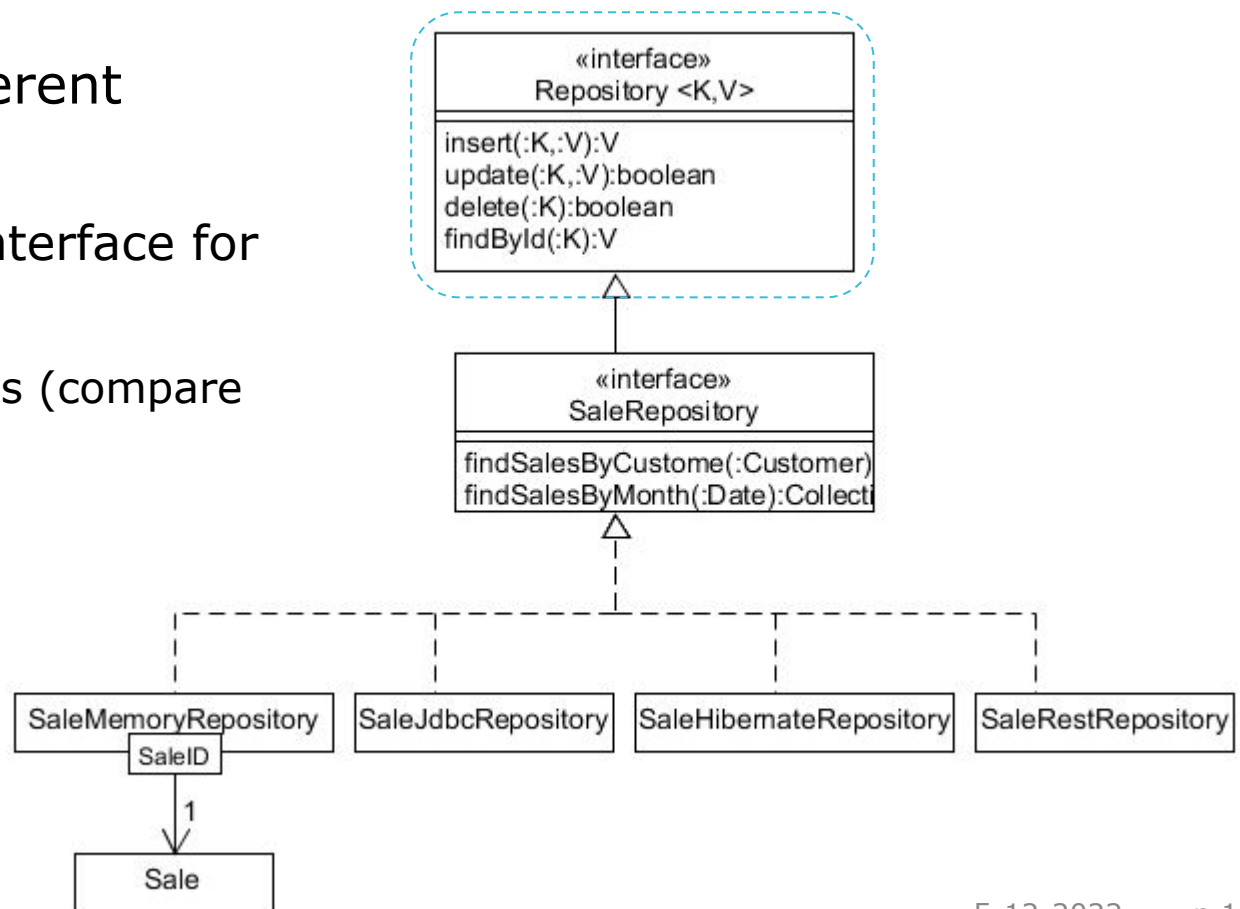
DDD: Repositories

- Repositories hide the differences between persistence technologies
 - Not always perfect: technology bleed



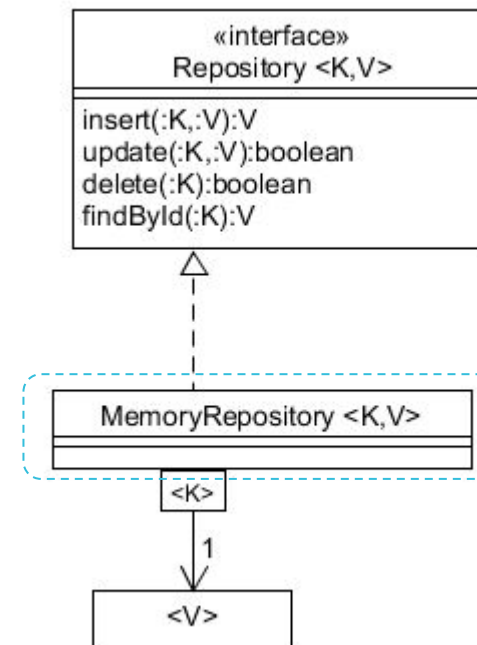
Generic Repository interface

- Different repositories do similar actions on one object (CRUD)
- Actions have different parameter types
 - Can you use an interface for this?
 - Yes, using generics (compare with Collections)



Generic Repository implementation

- Depending on the underlying technology a generic implementation can be feasible



Generic Repository implementation

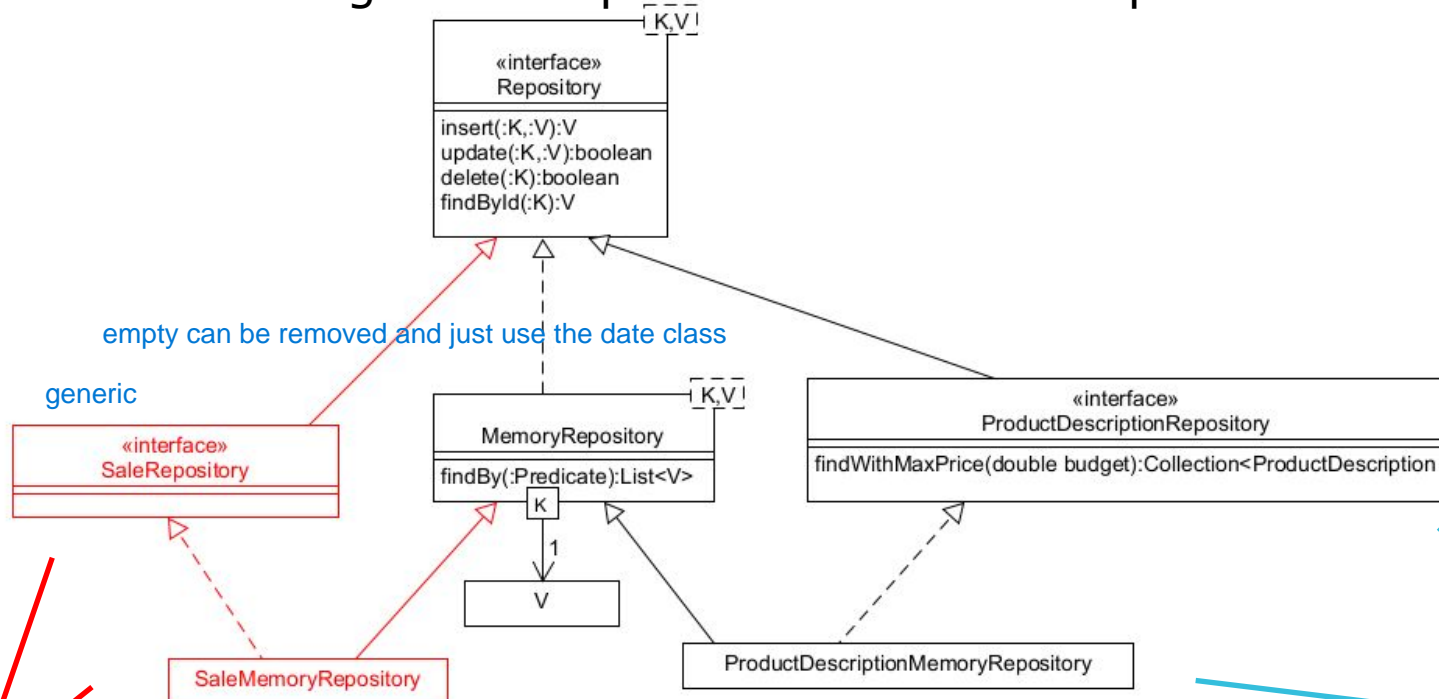
```
public class MemoryRepository<K,V> implements Repository<K, V> {  
    protected Map<K, V> data =new ConcurrentHashMap<>();  
  
    @Override  
    public boolean update(K key, V value) {  
        return data.put(key, value)!=null;  
    }  
  
    @Override  
    public V findById(K id) {  
        return data.get(id);  
    }  
  
    // + Some methods specific to Memory Repository  
    public List<V> findBy(Predicate<V> predicate) {  
        return findStreamBy(predicate).collect(Collectors.toList());  
    }  
  
    private Stream<V> findStreamBy(Predicate<V> predicate) {  
        return data.values().stream().filter(predicate);  
    }  
    //...  
}
```


Use of generic repository implementation

```
public class SaleManager {  
    private final Repository<Long,Sale> saleRepository =  
    new MemoryRepository<>();  
    ...  
}
```

Specific Repository Implementation

- Read methods often return different object selections
 - Collection `<V> findByXxx`
 - Subclass the generic implementation to add specific methods



findWithMaxPrice is not part of **MemoryRepository**
Specific subclass that adds this method.

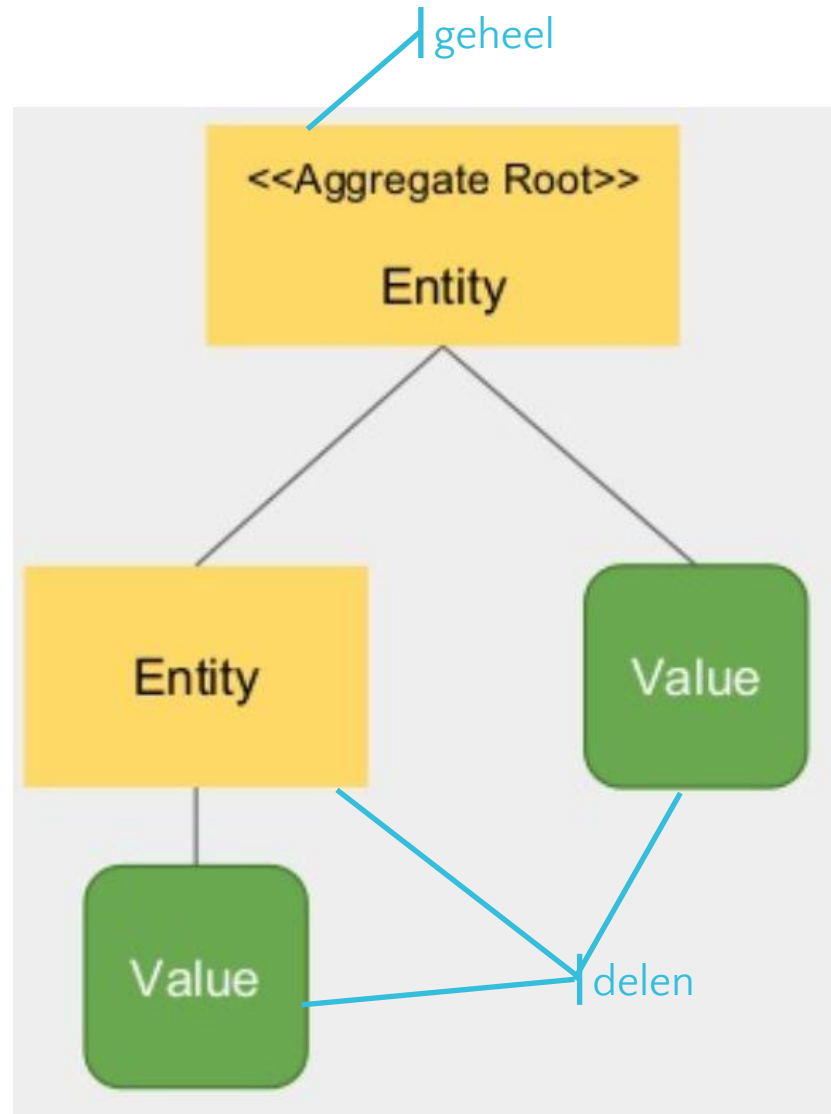
Specific Repository Implementation

```
// ProductRepository is called ProductCatalog in larman  
public interface ProductDescriptionRepository  
    extends Repository<Long,ProductDescription> {  
    Collection<ProductDescription> findWithMaxPrice(double budget);  
}
```

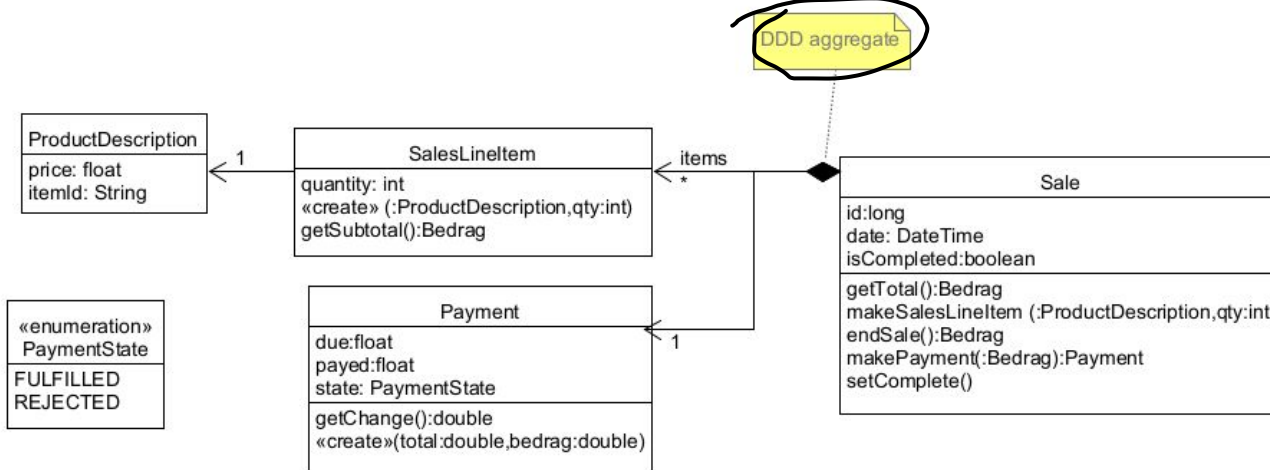
```
public class ProductDescriptionMemoryRepository  
    extends MemoryRepository <Long,ProductDescription>  
    implements ProductDescriptionRepository {  
  
    @Override  
    public Collection<ProductDescription> findWithMaxPrice(Money budget) {  
        return findBy(p -> p.getPrice().lowerThenOrEqual (budget));  
    }  
  
}
```

DDD aggregates

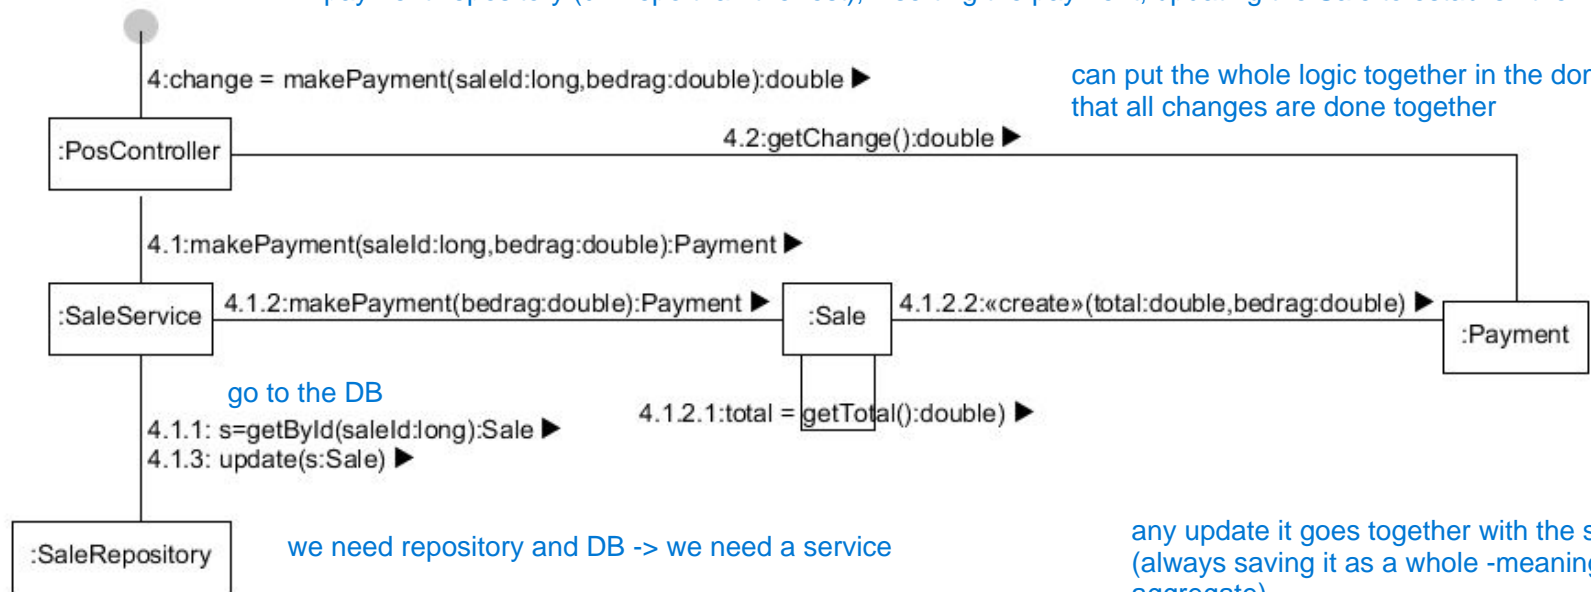
- All operations on the aggregate are always performed via the root
 - Business logic within aggregate implemented in domain, not in services
- Also for persistence => guideline: 1 repository per aggregate
 - All parts of the aggregate are retrieved and saved together



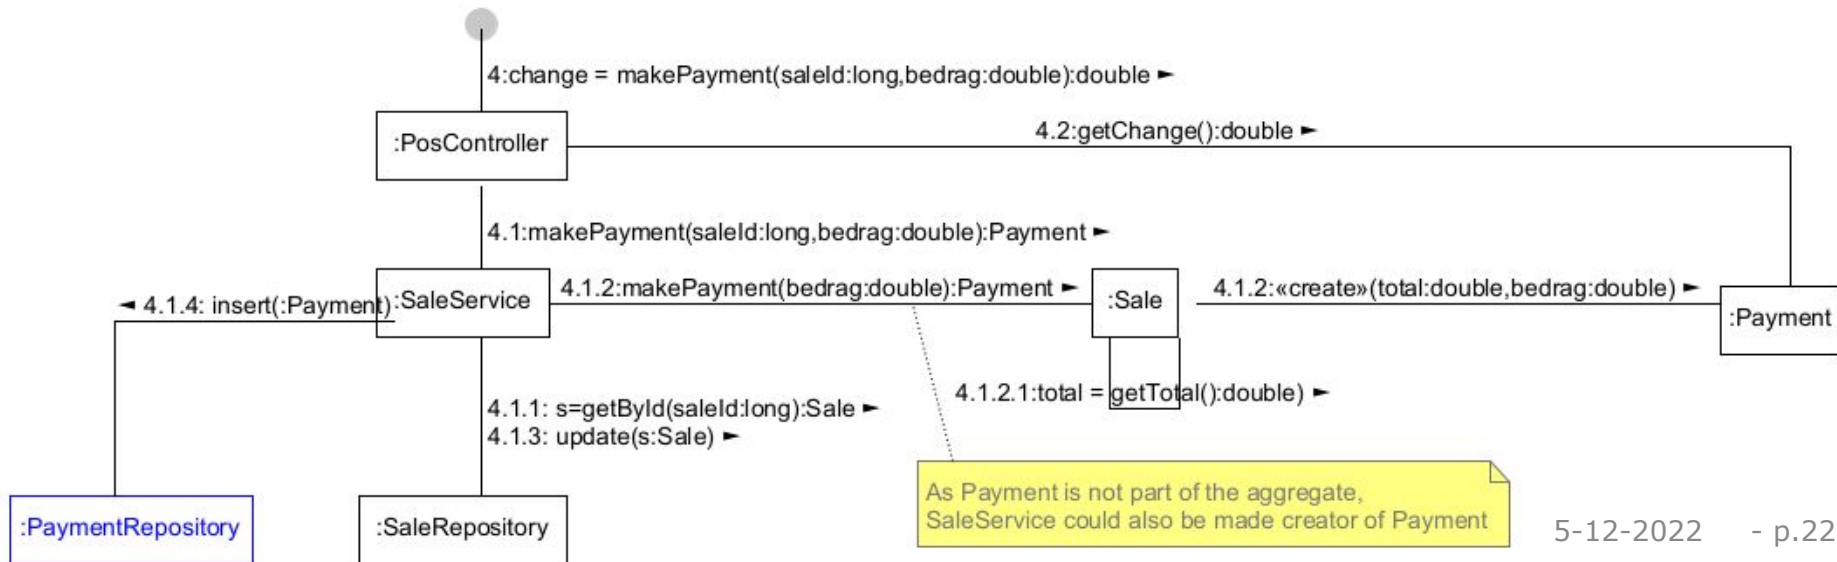
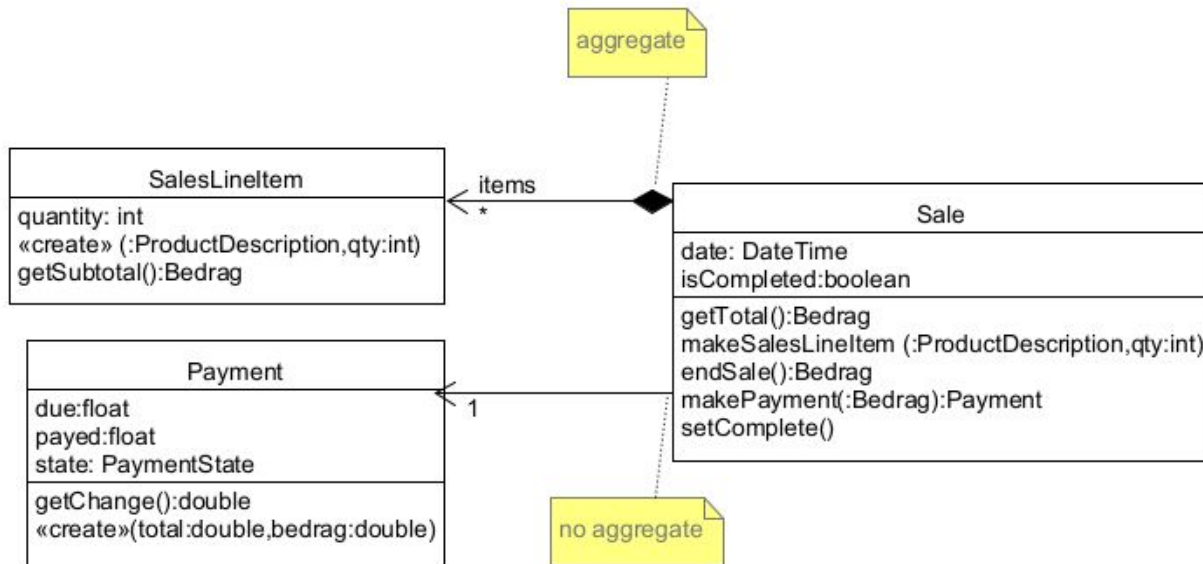
Example: Payment in Sale aggregate



if we take the payment out of the aggregate (remove rel with Sale) -> save the payment through a paymentRepository (diff repo than the rest), inserting the payment, updating the Sale to establish the relation



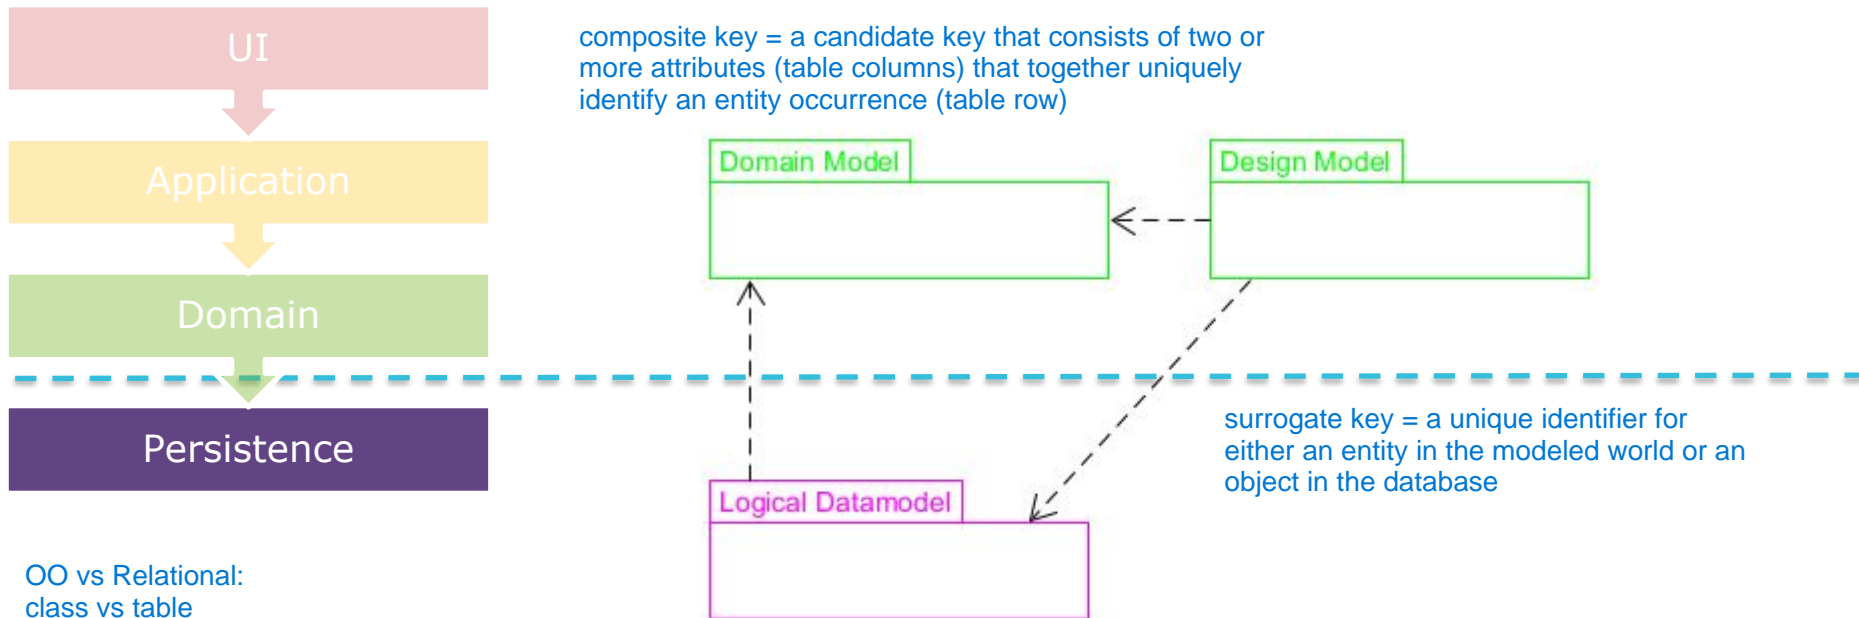
Example: Payment not in Sale aggregate



As Payment is not part of the aggregate,
SaleService could also be made creator of Payment

Datamodels

Designmodel and datamodel



composite key = a candidate key that consists of two or more attributes (table columns) that together uniquely identify an entity occurrence (table row)

surrogate key = a unique identifier for either an entity in the modeled world or an object in the database

OO vs Relational:

class vs table

associations 1 to many (natural relation in relational, done through a List in OO)

1 to 1 (embedded in relational)

many to many - can't be done in the relational world - it is done using a 1 to many using an FK

attribute matches to a column

object maps to a row

behavior centric (partition is based on behavior) vs data centric (partition is based on data -that's why it's embedded)

associations are based on pointers (memory reference) associations are based on keys (based on FKs)

ID is the memory location vs ID is the PK

natural relation with a pointer = 1 to 1

pointer points in one direction (relations are unidirectional) vs bidirectional (can use the keys anywhere)

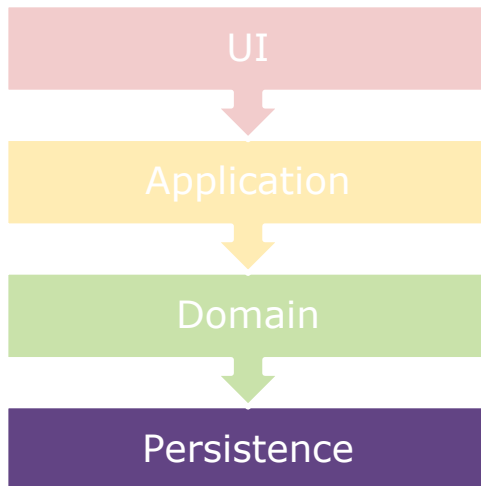
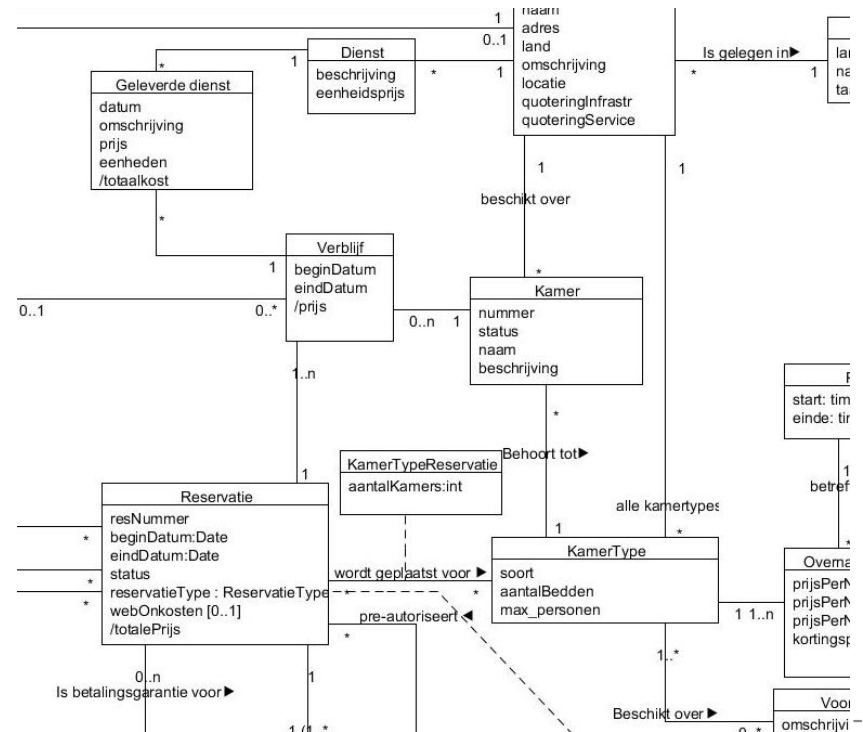

2 relations to make it bidirectional

memory vs persistent

methods vs there are none (closest thing to methods are stored procedures/trigger)

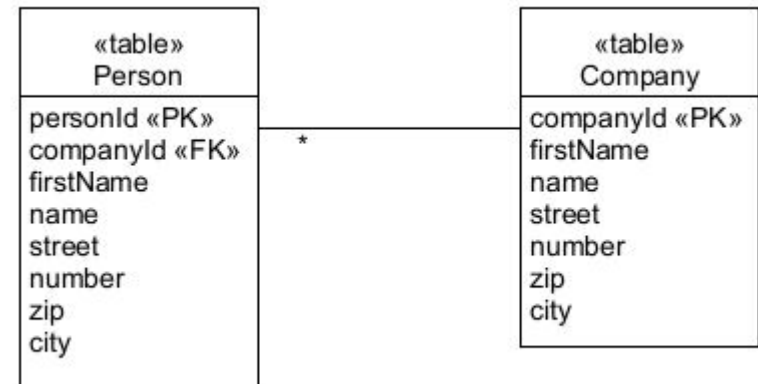
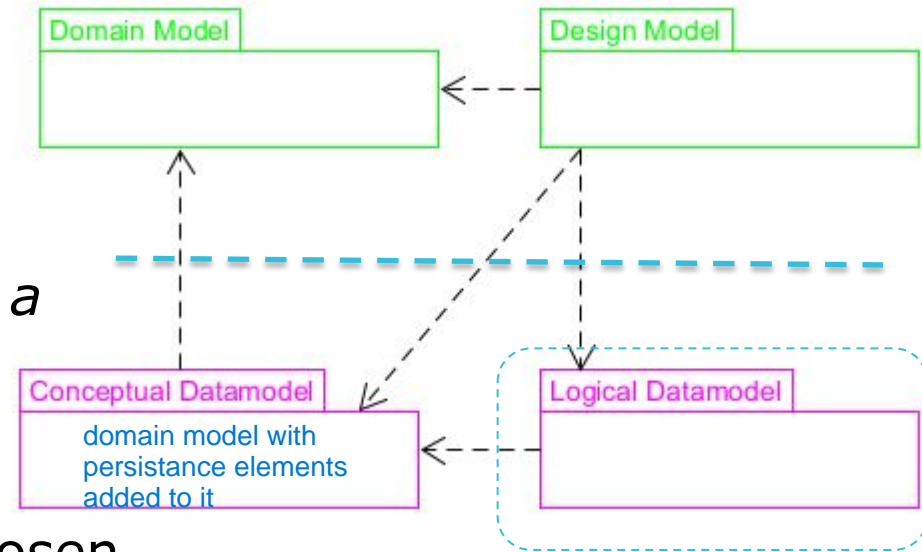
Domain model

- ❑ Class diagram with conceptual classes
- ❑ Part of analysis



Logical *Datamodel*


- ❑ Only elements that are persisted
(most classes, but not e.g. a pure behavioural class like *SalePricingStrategy*: does not need to be persisted)
- ❑ Storage technology was chosen
- ❑ Logical structure of datastructure
- ❑ Conventions
 - ❑ Use stereotypes for relational datamodel
 - ❑ «table»
 - ❑ «PK»
 - ❑ «FK»
 - ❑ Table names: plural



Note: sometimes people differentiate between **conceptual** datamodel (technology agnostic) and **logical** datamodel (technology specific)

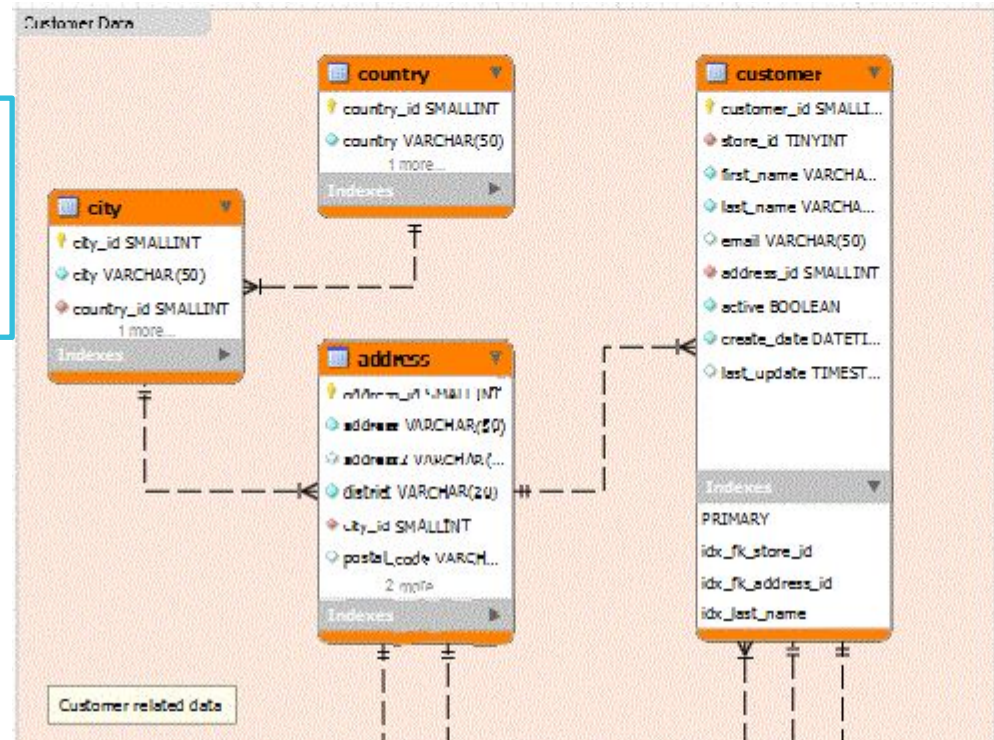
Physical datamodel

- ❑ Physical structure
database, product
specific
- ❑ Exact data types
- ❑ Implementation aspects




Often built using a database tool

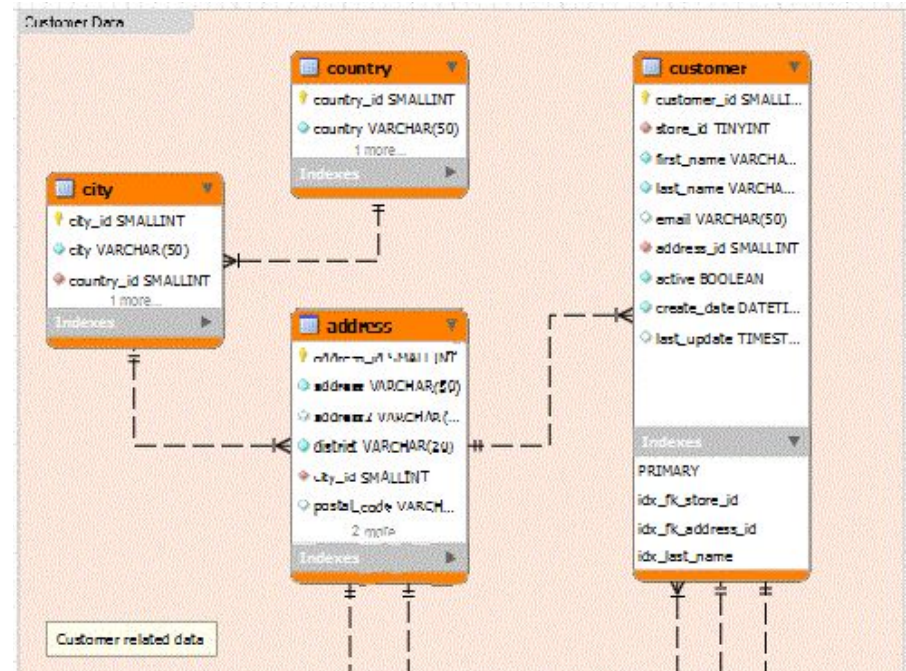
*e.g.. SQL Developer Data Modeler ,
IntelliJ, mySQL workbench...*




Physical datamodel



- ❑ Implementation aspects are often performance related
 - ❑ Define indexes on columns that are frequently searched
 - ❑ Denormalise
 - ❑ Product specific features:
 - ❑ Partitioning tables
 - ❑ Clustering tables
 - ❑ Special indexes: bitmap index

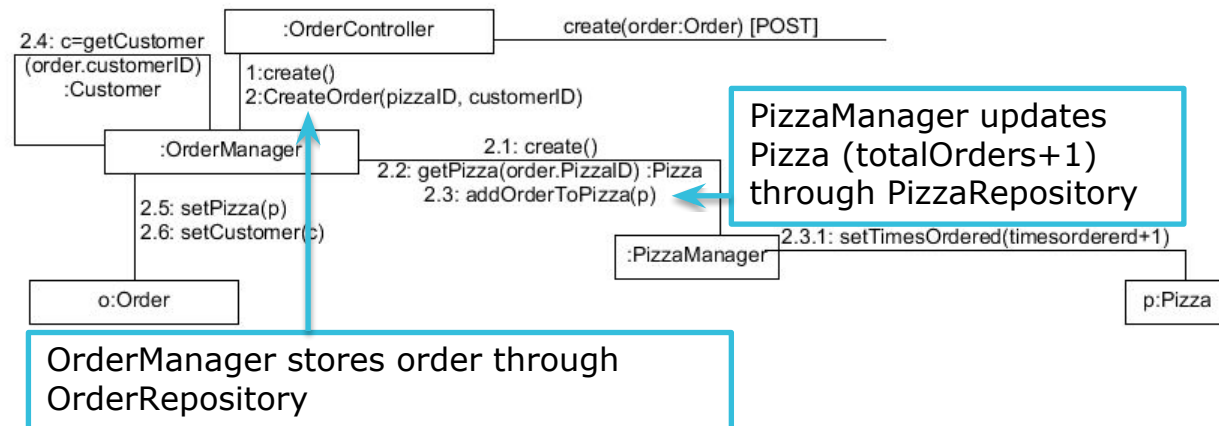


Transactions



Transactions

- Transaction demarcation is a business decision, typically at the level of a technology agnostic service
- A service can invoke operations on multiple repositories
- Just like for databases, transactions should fail or succeed as a whole

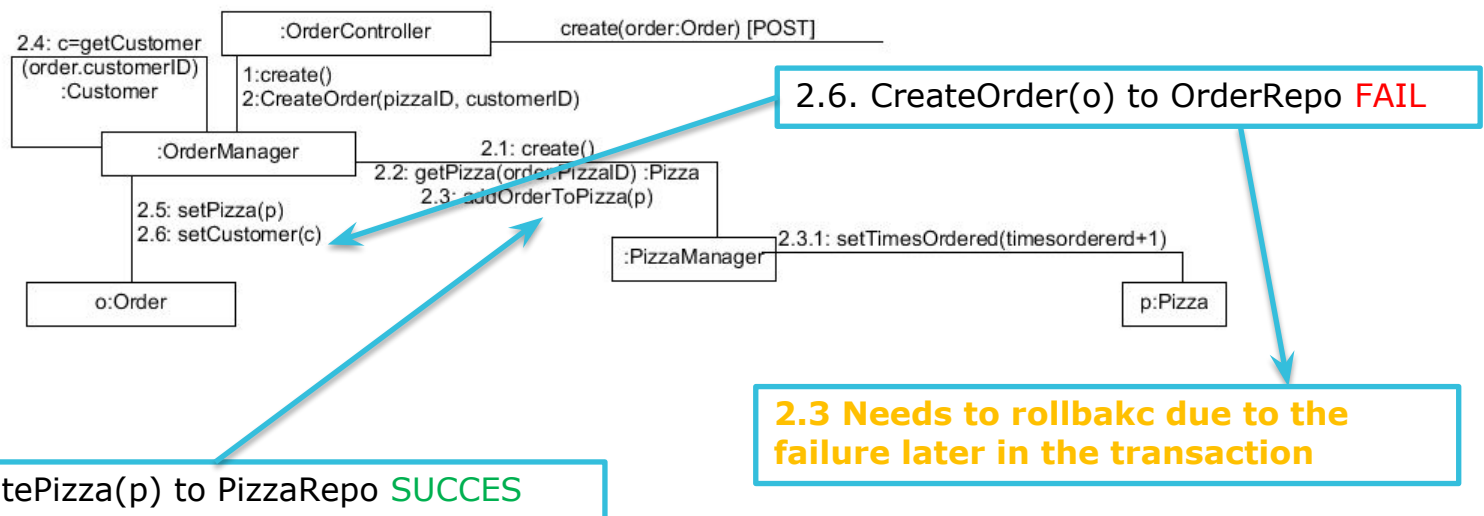


Transactions

- Transaction demarcation is a business decision, typically at the level of a technology agnostic service
- A service can invoke operations on multiple repositories.
- Just like for databases, **transactions** should fail or succeed **as a whole**

transaction ends when the method ends

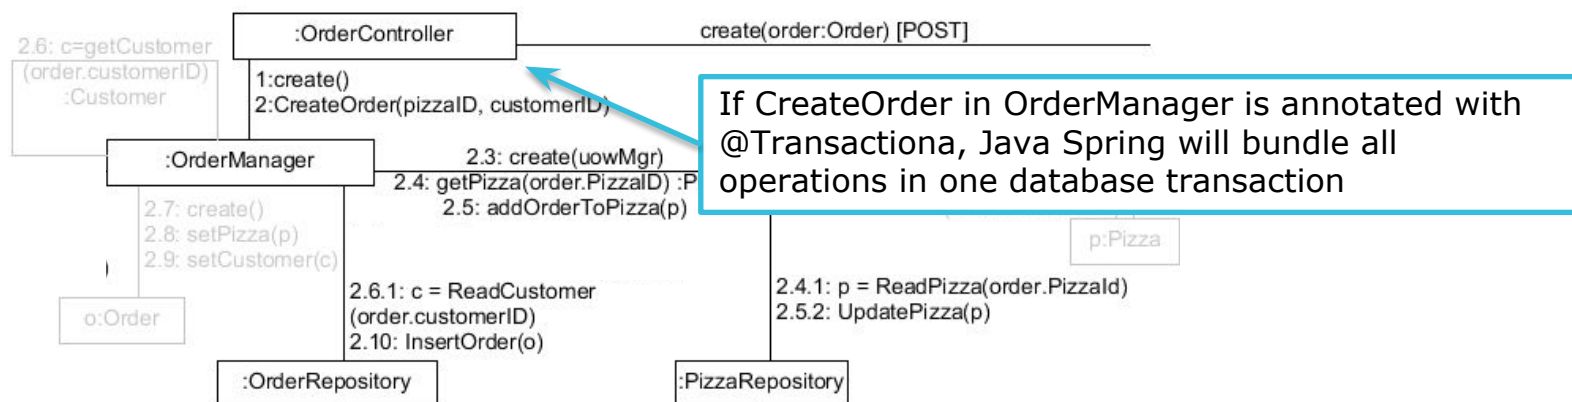
independent from DB



Transaction implementations

- Libraries offer solutions to implement transactions
- Java example
 - In Spring you can annotate a method with `@Transactional`
 - Before the first database modification Spring automatically starts a transaction
 - If all goes well the transaction is committed at the end of the method
 - If the method throws an exception the transaction is rolled back

can modify behaviors



Object-Relational Mapping

A hand is holding a light-colored wooden block, positioned as if about to place it into a circular hole in a wooden surface. The scene is lit with warm, directional light from the left, creating strong shadows. On the right side of the image, there is a large, thick, black stylized letter 'C' that partially obscures the background.

Object

- Class
- Attribute
- Object
- Association
- Method

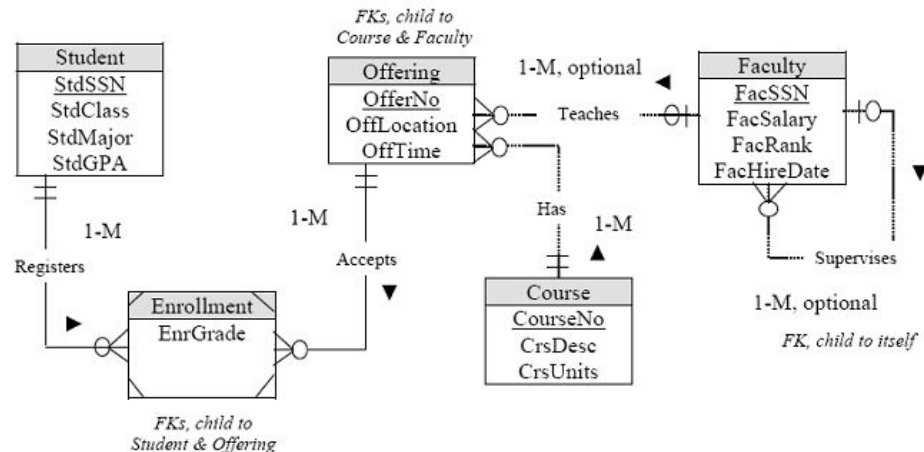
Relational

- Table
- Column
- Row
- Relation
- Stored Procedure

Many Object Oriented concepts have a natural relational counterpart But there are differences as well...


OR gap: Persistent \Leftrightarrow Memory

- relational database is **persistent**
 - Entity associations use a key
 - = additional persistent attribute
 - A row contains one foreign key
 - A primary key can be referenced by many rows
 - Simplest relational association is a bidirectional 1-* association
 - *-* association: not supported. Split up in two 1-* associations.



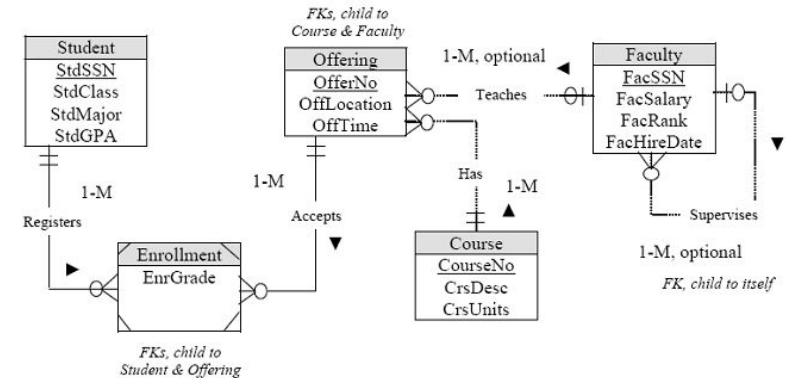
OR gap: Persistent \Leftrightarrow Memory

- A (Object Oriented) program is **transient**
 - Entity associations use a (memory) reference
 - References one thing
 - The same memory location can be referenced by many entities
 - Simplest memory relation is a unidirectional $* \rightarrow 1$ relation
 - $\rightarrow *$ (... to Many) association: refer to a Collection
 - Bidirectional association: additional unidirectional association in the opposite direction
 - $1 \rightarrow$ (One to ...) association: do not expose objects you want exclusive access to (or add extra logic)




OR Gap: Data \Leftrightarrow Behaviour

- Relational: Data focus, based on mathematical model
 - stored procedures added later
- 1-1 relations are few
- SQL is a declarative language
 - Specify WHAT you want
- DBMS determines steps (plan) to manage data



OR Gap: Data \Leftrightarrow Behaviour

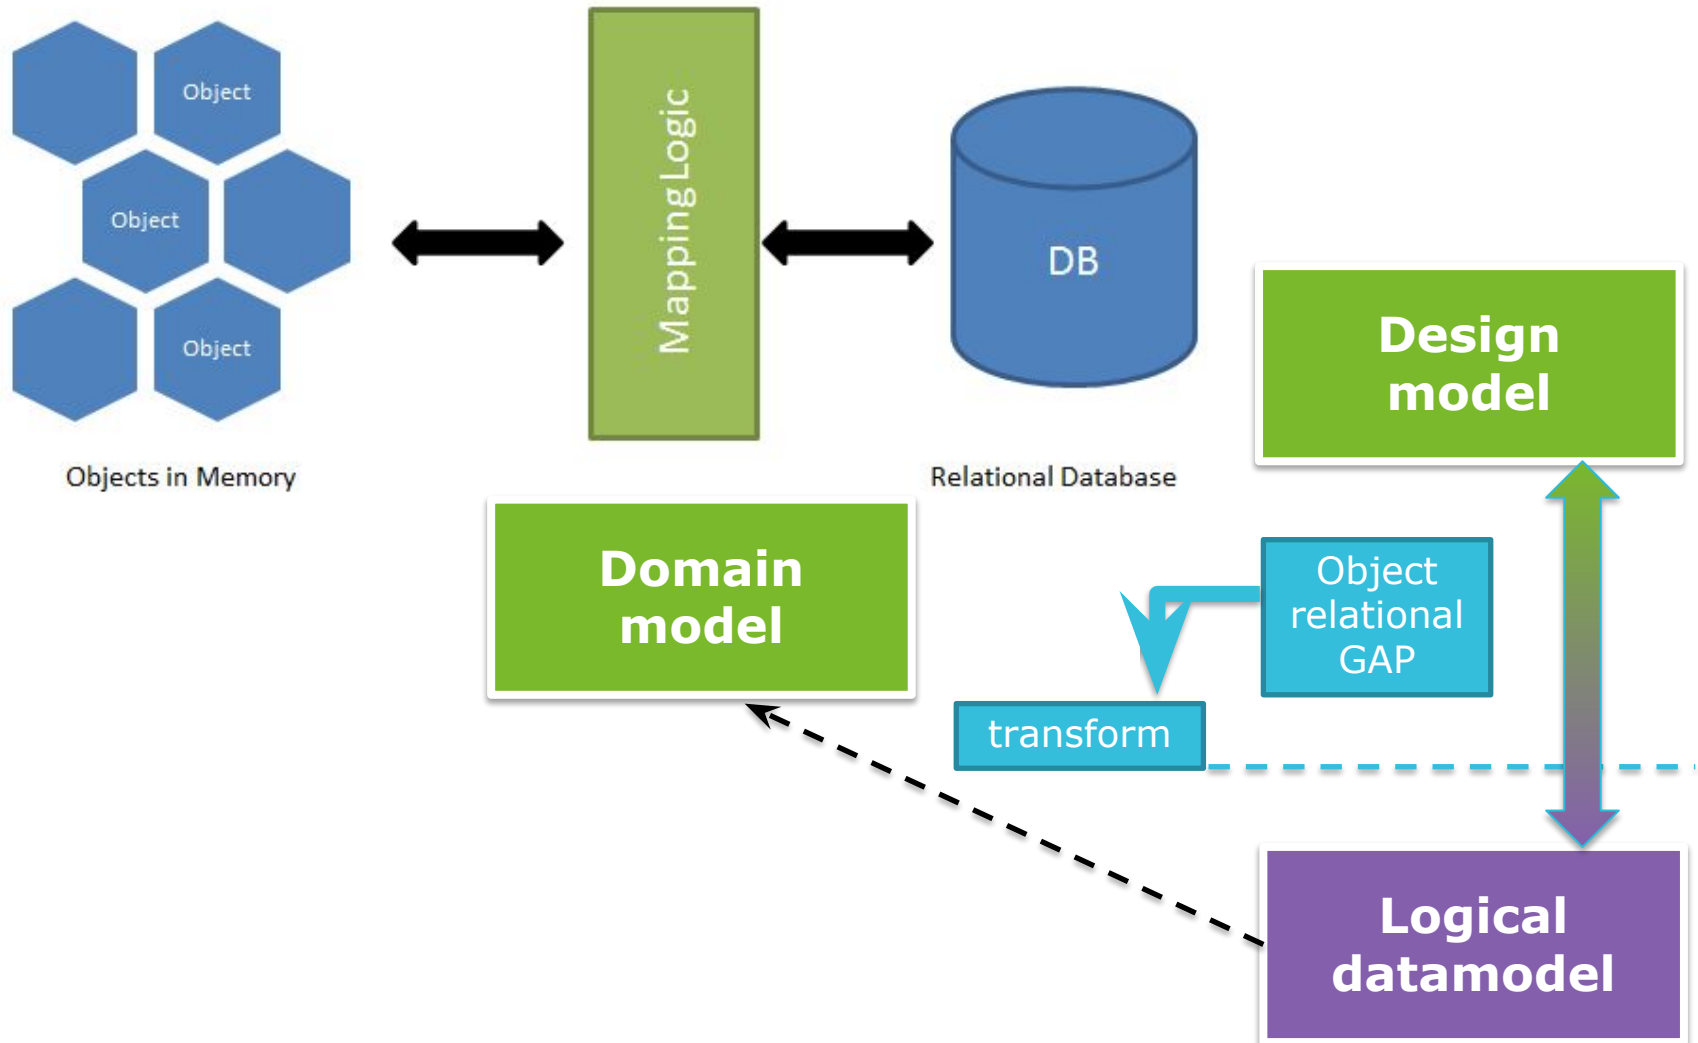
- Object oriented: combines data and behaviour
 - encapsulation and data hiding
 - An attributes with its own behaviour gets a separate class, referenced in a \rightarrow 1 relation
 - methods are imperative
 - Specify HOW your want to do something
 - has steps, loops...
- Extra association type: inheritance
 - polymorphism, abstraction



How to bridge the Object-Relational gap?

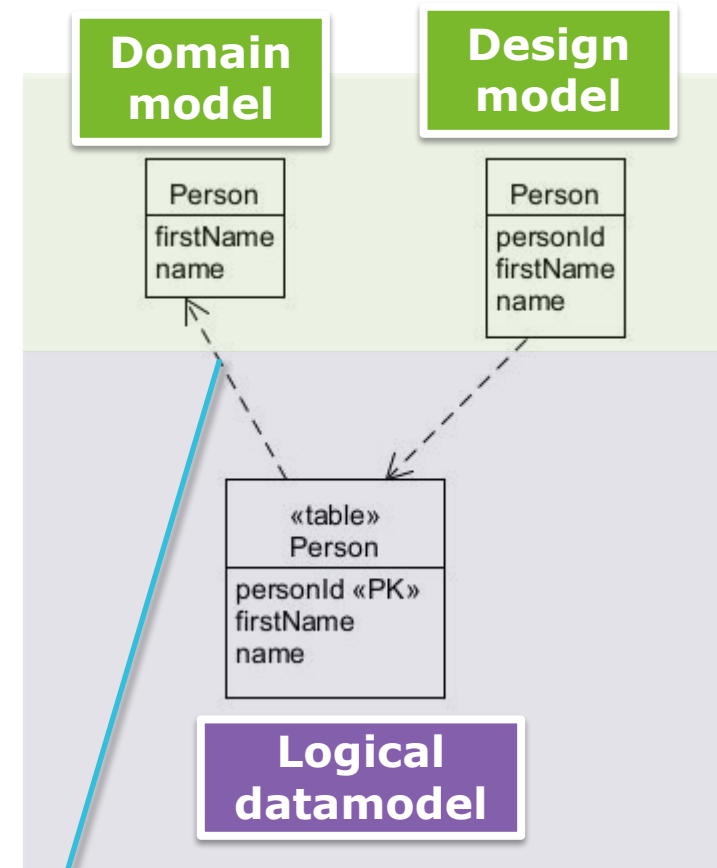
Object – Relational Mapping

O/R Mapping



Object identification

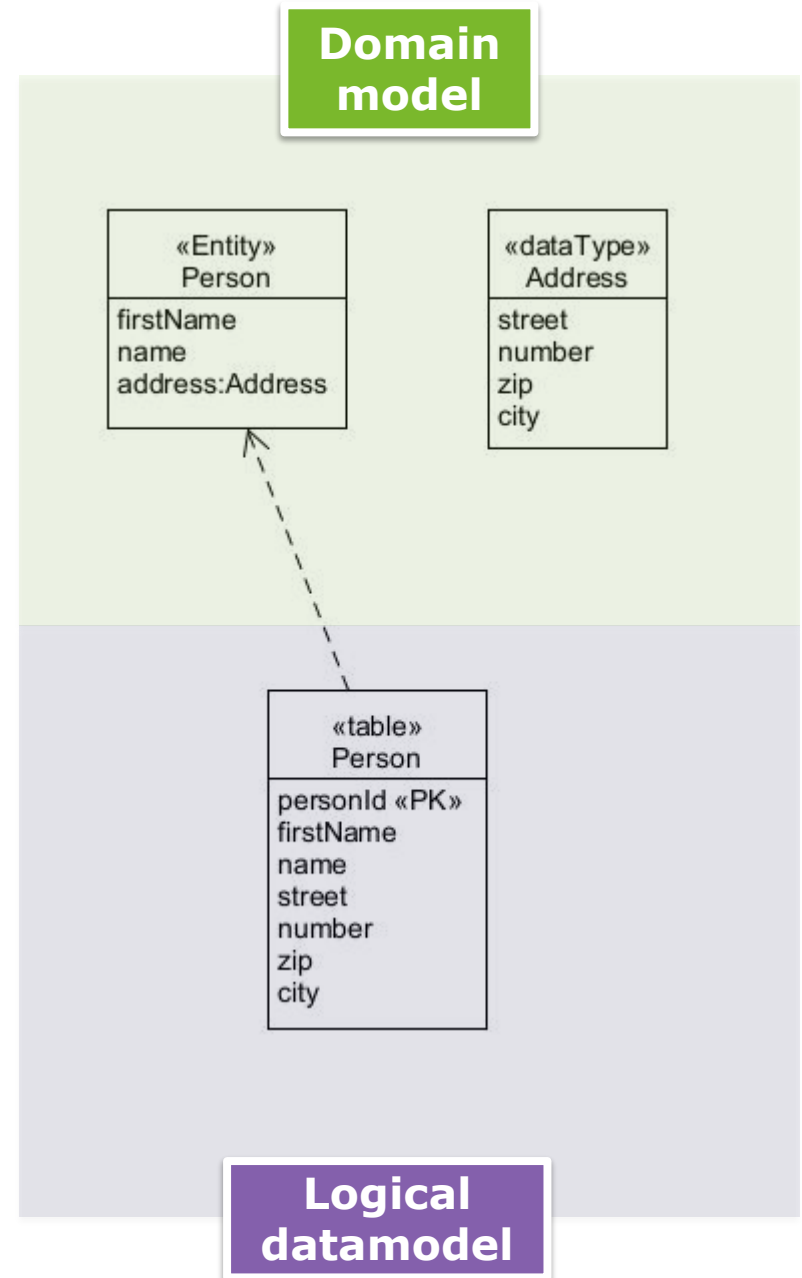
- ❑ OO Objects identified by memory location
- ❑ RD Rows identified by key
- ❑ OO -> RD requires adding identity (PK)
 - ❑ Surrogate key (artificial identity)
 - OR
 - ❑ One or more existing attributes
- ❑ If a surrogate key is added in the RD, it needs to be added to the OO design model as well



Note: the mapping is done in the direction opposite to the dependencies

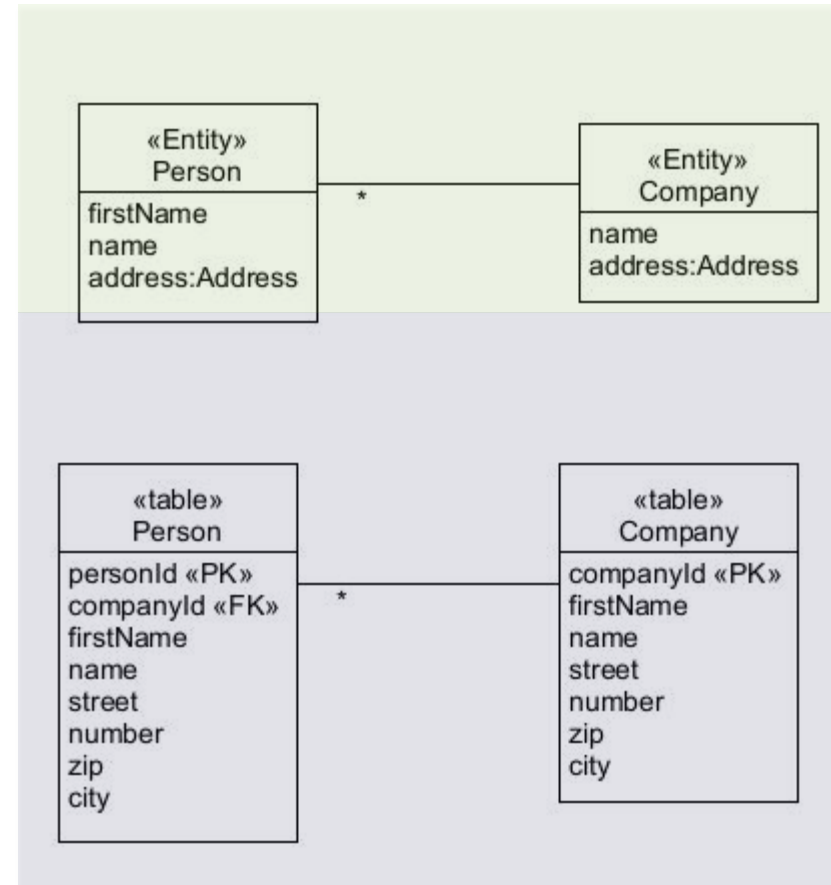
Mapping relations: 1- 1

- Add the attributes of the Value Object to the Entity table
- 1-1 DDD aggregates can also be mapped like this



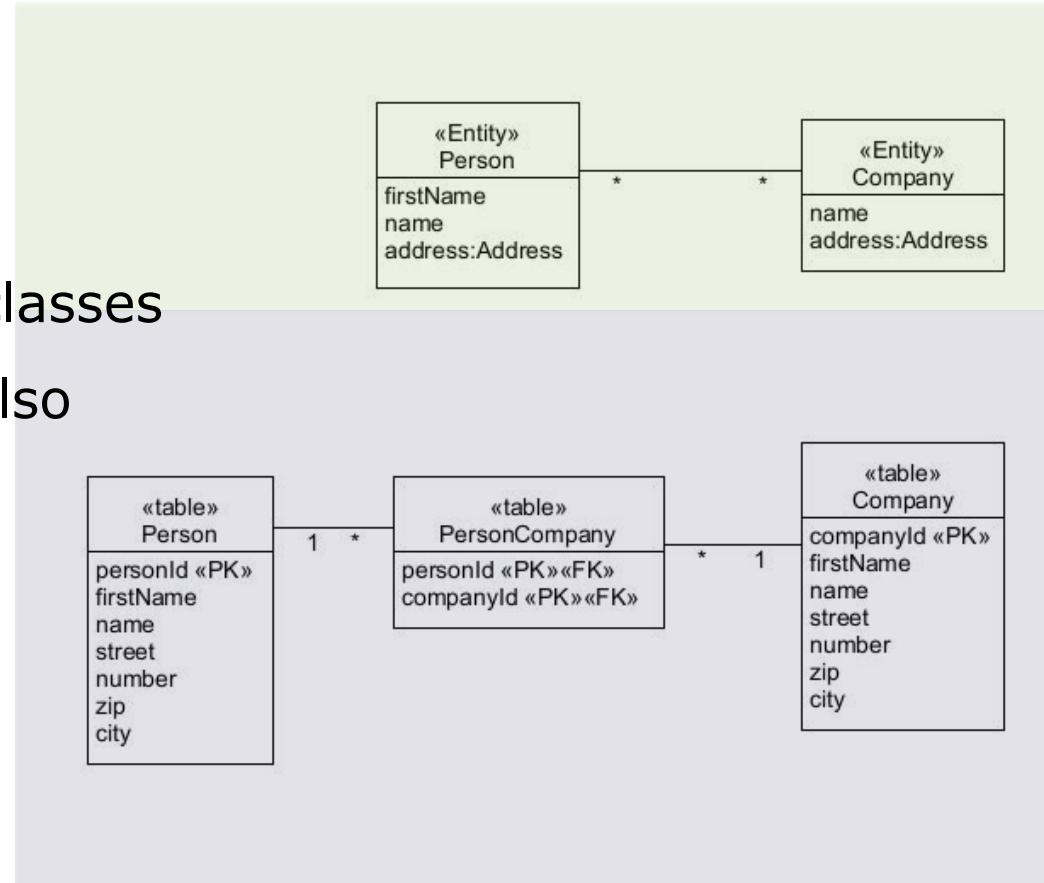
Mapping relations: 1- *

- ❑ Each class maps to a table
- ❑ FK goes to the *-side



Mapping relations: *-*

- ❑ In relational model transform
- to 2 1-* relations
- ❑ Extra table
(PersonCompany)
has FK for both domain classes
- ❑ Association classes are also
mapped in this way



Inheritance mapping strategies

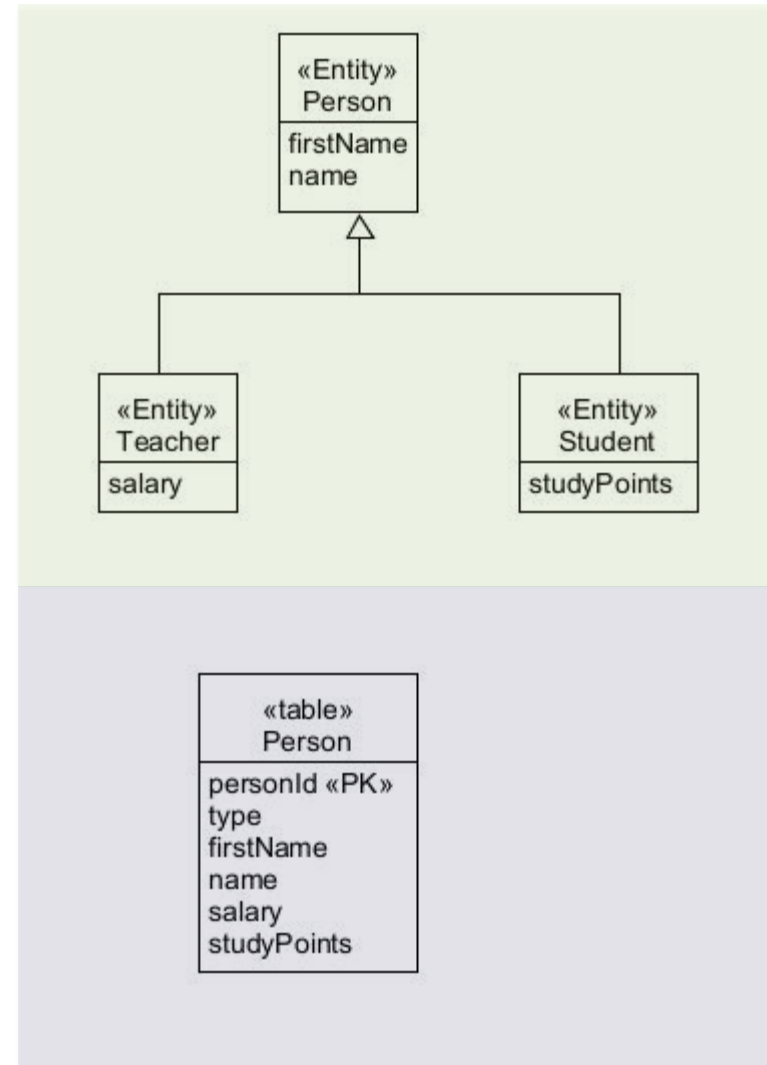
1. Table / hierarchy (1 table)

- ❑ Add type attribute
- ❑ All subclass attributes are optional
 - ❑ No NOT NULL constraints for subclass attributes
 - ❑ Empty fields
 - ❑ Large, complex table

Views can mitigate this

2. Table / subclass (2 tables)

3. Table / class (3 tables)



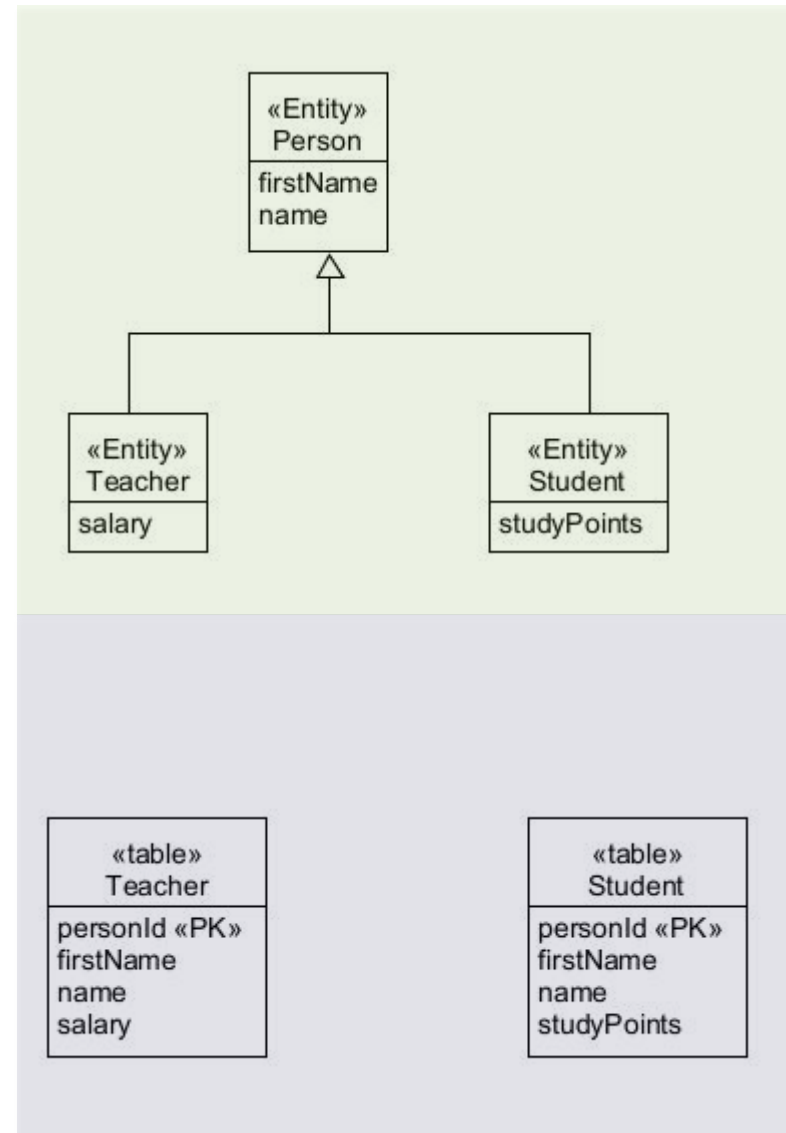
Inheritance mapping strategies

1. Table / hierarchy (1 table)

2. Table/subclass (2 tables)

- superclass query requires UNION
- Superclass relations need to be implemented for each subclass

3. Table / class (3 tables)



Inheritance mapping strategies

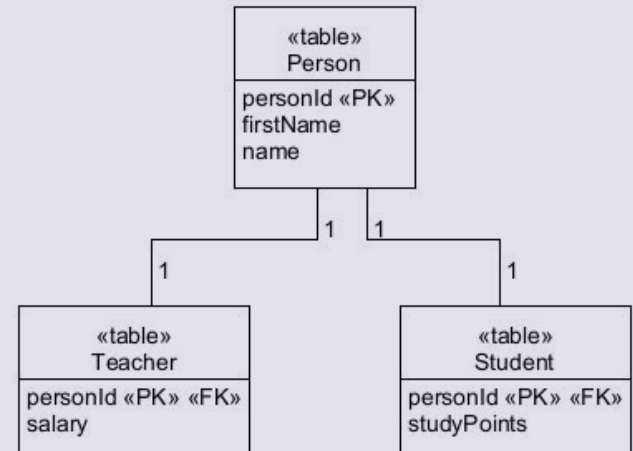
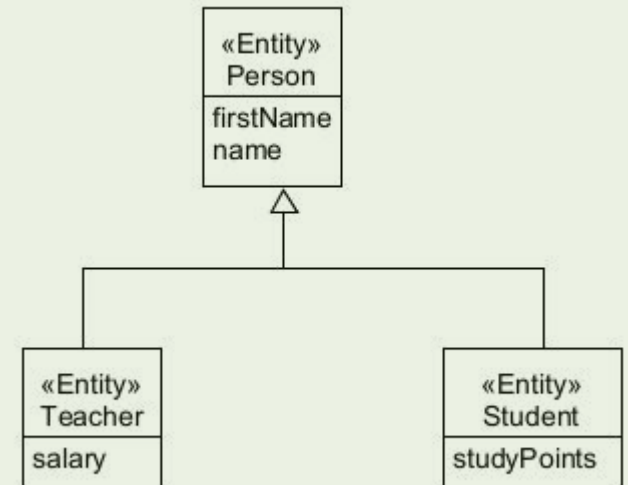
1. Table / hierarchy (1 table)

2. Table/subclass (2 tables)

3. Table / class (3 tables)

- ✓ extra subclasses can be added without modifying existing tables
- ✓ "multiple inheritance": a teacher can also be a student

□ JOIN needed



Inheritance mapping strategies

1. Table / hierarchy (1 table)

- Fastest for reading (**most popular**)

2. Table / subclass (2 tables)

- Useful if operations across subclasses are uncommon (example infrequent polymorphic queries)



3. Table / class (3 tables)

- Useful if database structure changes often (**least popular**)

Object

- Transient
- Class
 - Attribute
 - Inheritance
- Object
 - instance attribute
 - implicit ID: reference
- Relation
 - reference
 - Unidirectional
 - Ordered (List, array...)

Relational

- Persistent  Red: difference
- Table  Green: match
 - Column
- Row
 - field
 - explicit ID: primary key
- Relation
 - foreign key
 - bidirectional

Object

- Behaviour focus
 - imperative (how)
 - Works with objects
- Methods
 - interfaces
 - encapsulation
 - Data hiding

Relational

- Data focus
 - Declarative (what)
 - works with sets
- Stored procedures*
- Triggers, constraints
 - Permissions

Some similarity

* Later addition

Summary



1. Persistence stores
2. Repositories
3. Datamodels
4. Transactions
5. Object-Relational Mapping